

# PDF test fixtures

The PDF integration tests (`PdfIntegrationTests`) look for the files below. **They are skipped (not failed)** if a file is missing, so the suite stays green until you add them.

Create real PDFs (any tool — Word "Save as PDF", LibreOffice, an online generator, or print to PDF). Use a **clear, printed font  $\geq 18$  pt** (these are OCR'd, so avoid tiny or decorative fonts). Born-digital text PDFs are fine — they're rasterized and OCR'd like a scan.

Put the files **directly in this folder** (`test/assets/pdf/`). They're copied to the test output and discovered automatically.

---

## 1. `simple_en.pdf` — **exactly 1 page**

Must contain this text, clearly printed:

The quick brown fox jumps over the lazy dog  
1234567890

Tested: single-page extraction, digit allowlist, and searchable-PDF round-trip. Assertions look for the words **QUICK**, **BROWN** and the number **1234567890**.

---

## 2. `multipage_en.pdf` — **exactly 3 pages**

One distinct word per page, large and centered:

| Page | Text    |
|------|---------|
| 1    | ALPHA   |
| 2    | BRAVO   |
| 3    | CHARLIE |

Tested: multi-page extraction and **page ordering** (page 1 → ALPHA, page 2 → BRAVO, page 3 → CHARLIE).

---

## 3. `scanned_doc.pdf` — **a real scan (optional, recommended)**

A genuinely scanned or photographed document saved as PDF (image-only pages — the realistic OCR case). Any content; ideally a page or two of clear printed text.

Tested by `ScannedPdfTests.Real_scanned_pdfextracts_meaningful_text`, which just asserts that meaningful text comes out ( $\geq 10$  characters). Tip: phone-scan a printed page or "Print to PDF" a scan.

*Note: scanned-image PDFs are **also tested automatically** — `ScannedPdfTests` builds an image-only ("scanned") PDF from `sample.png` in-process and OCRs it, so the scanned path is covered even without this file.*

## Optional extra image fixtures

For even more realism you can drop images into `test/assets/images/` (receipts, license plates, signs, photos at an angle, faint scans). They're copied to the test output automatically; add matching tests as needed. Non-Latin scripts work for *extraction*, but the searchable-PDF text layer is optimized for Latin text.